

In such cases it is crucial to be on the line with the customer just to let them know that the real person is working on their case. So before getting back to John I would quickly (several minutes) go through the admin dashboard and their service to gather initial information, if the information John provided is real in the DB or if it's just a front-end glitch, and check their service if it could be restarted or brought back online quickly.

If the issue is not that easy and could not be resolved on the spot I will get back to John. The email will be sent in several minutes timeframe after the quick check explained higher:

Hello John,

Thank you for contacting the RunPod's support team.

I understand how critical your deadline is so I just wanted to get back to you to inform you that I have personally taken ownership of your case, and I will be keeping you posted on any insights I discover.

If you think that there's any piece of information that may help with understanding the issue deeper please do not hesitate to provide it. If any additional information from your side is required I will let you know right away.

Thank you.

Best regards, Vitalii.

1. There may be a lot of issues around this like:
 - a. Front-end issue, the dashboard glitched and shows wrong information about the billing status and the status of the service;
 - b. Network issue that prohibits John from connecting to his pod, disconnection on our side, network disruption, connectivity issues;
 - c. Hardware issue, the physical GPU could be broken or cracked;
 - d. Billing issue, John could overuse the resources and due to the zero balance the service shut down, and the low balance reminder email didn't go through;
 - e. Software issue, John's model could actually run into an issue and brake something
 - f. API server issue, it may be down, or it may incorrectly address the calls;
2. Taking into account that John is currently stressed, frustrated, and probably even angry I will keep the empathic and non-blaming tone, using passive voice and emphasizing that the team is working to find the root cause of the issue on our end and that we need help from John's side as well. However, if it's an issue on our side I will be truthful, and clear.

a. Hello John

Thank you for your patience.

While the team is investigating this case from our end we would require some help from your side. Please provide the following information that may help us understand the issue quicker:

- When was the last time you were able to log in to your instance?
- Do you remember anything strange happening during your last log in?
- Did you start any new processes while running the main one?
-Any additional information that I may need to understand the issue

I strongly appreciate your help here and will be in touch as soon as I have any updates.

- b. If I'm not able in a reasonable time to identify the issue and help the customer using all the knowledge I have, the escalation will be considered. The internal escalation guidelines should be used for this process. But briefly, I would provide a summary of the incident with the most important facts about the incident, like identification information, incident timelines information, logs, and other technical info I consider important for the engineering team.
3. After the initial look I would check everything else. Check the DB to understand the real state of the service, check the monitoring systems to understand the status of the pod. Check if it can be brought back without major impact, check if the data can be restored. Check API logs, check the logs of the container. Check if there are no other billing-related issues like this;
4. Using the company's guidelines, the incident-ticket must be created including all the information engineering team requires. The proper priority must be set. The proper point of contact must be notified right away.
5. I would do almost the same at the beginning, putting all the effort into investigating, understanding and resolving the issue using all my knowledge and troubleshooting skills. However, if I run into a wall, I will still be trying to raise the question with the team and have someone take this case. Due to the type of services the company provides there always should be an on-duty engineer in case of a system outage or major clients issues. Yes, it may take a little more time to bring someone in, or it may take more than one try, but again, I cannot just leave this customer (high paying customer) hanging there with their problem, I need to make sure that their issue is resolved no matter what.
6. If we are seeing that fixing the issue may take longer than expected we may want to make sure that John's work can be continued, ideally from the place he left everything, restore the data and deploy it to another instance, give John more resources to be able to finish this project within the deadline.
7. First of all the customer must confirm that their issue is resolved and they can continue working with no further interruption. Second I may need to confirm with the management granting John a compensation (balance bonus, free upgrade of their service) for this incident according to the company's policies. Third, we need to make sure that the root cause of this incident is fixed and will not happen again. Do a post-mortem of the incident, brainstorm the ideas of how to avoid such cases in the past, document everything and post the most important insights and findings into the internal KB.