

It was very interesting to get to know the platform and go through this task.

In the zip attached you will find the following files:

- Dockerfile
- [handler.py](#)
- requirements.txt
- [runpodApiRequest.py](#) (the file to send an API request to the instance)
- screenshot of the prompt and generated result

At first I was trying to do everything on my machine locally download the model create a build and then deploy it to the pod, so the model download would only happen once, but I realized that I did not have enough RAM and had to create a build without the model in it.

Additionally, after reading the documentation I found that there's a more efficient way of using huggingface models through the cached models:)